

Full Pipeline Using Spark

Ashwanth Samuel (as7cs)

MSDS Class of '19

Motivation

- This dataset describes songs from Spotify in 2017 with attributes utilizing the streaming services' API.
- There are 16 attributes including acoustiness, danceability, duration in milliseconds, energy, loudness etc.
- I wanted to run a simple linear regression using Spark and MLlib software and that is detailed here in these slides.

I ultimately decided that I wanted to try and predict the energy of a song from its “speechiness” factor.

Creation of Context

```
In [1]: import pyspark
conf = pyspark.SparkConf().setAppName('odl').setMaster('local')
sc = pyspark.SparkContext(conf=conf)
sqlc = pyspark.sql.SQLContext(sc)
```

```
In [2]: import pyspark.sql.functions as sf
```

Preparation of Data

```
In [3]: dataPath = "data.csv"
df = sqlc.read.format("csv")\
    .option("header", "true")\
    .option("inferSchema", "true")\
    .load(dataPath)
```

```
In [5]: parquetPath = '/home/ec2-user/SageMaker/Open-Data-Lab/education/as7cs'
df.write.parquet(parquetPath)
```

```
In [13]: # create train/test sets
seed = 7
(testDF, trainingDF) = final.randomSplit((0.20, 0.80), seed=seed)
print('training set N = {}, test set N = {}'.format(trainingDF.count(), testDF.count()))

training set N = 1605, test set N = 412
```

```
In [14]: from pyspark.ml.linalg import Vectors, VectorUDT
```

```
⌘ In [15]: # make a user defined function (udf)
sqlc.registerFunction("oneElementVec", lambda d: Vectors.dense([d]), returnType=VectorUDT())

# vectorize the data frames
trainingDF = trainingDF.selectExpr("speechiness", "oneElementVec(energy) as energy")
testDF = testDF.selectExpr("speechiness", "oneElementVec(energy) as energy")

print(testDF.orderBy(testDF.speechiness.desc()).limit(5))

DataFrame[speechiness: double, energy: vector]
```

```
In [16]: # rename to make ML engine happy
trainingDF = trainingDF.withColumnRenamed("speechiness", "label").withColumnRenamed("energy", "features")
testDF = testDF.withColumnRenamed("speechiness", "label").withColumnRenamed("energy", "features")
```

```
In [17]: from pyspark.ml.regression import LinearRegression, LinearRegressionModel

lr = LinearRegression()
lrModel = lr.fit(trainingDF)
```

```
In [18]: type(lrModel)
```

```
Out[18]: pyspark.ml.regression.LinearRegressionModel
```

```
⌘ In [19]: predictionsAndLabelsDF = lrModel.transform(testDF)



print(predictionsAndLabelsDF.orderBy(predictionsAndLabelsDF.label.desc()).take(5))

[Row(label=0.622, features=DenseVector([0.716]), prediction=0.09321535682780889), Row(label=0.542, features=DenseVect
, features=DenseVector([0.412]), prediction=0.0813207831
iction=0.0984974866011892), Row(label=0.447, features=D
```

Code

Visuals

The right graph is a plot of energy versus danceability. From this we can see that the model is doomed to fail from the start. The bottom visual is a scatter plot of energy and danceability, which shows a little more promise for future analysis.



```
In [10]: print("Pearson's r(danceability,duration_ms) = {}".format(df.corr("danceability", "duration_ms")))
print("Pearson's r(danceability,speechiness) = {}".format(df.corr("danceability", "speechiness")))
print("Pearson's r(speechiness,energy) = {}".format(df.corr("speechiness", "energy")))
```

```
Pearson's r(danceability,duration_ms) = 0.004695083530137932
Pearson's r(danceability,speechiness) = 0.14266052295556916
Pearson's r(speechiness,energy) = 0.09310231843447833
```